



# Autonomous Systems: Deep Learning

## Deep Reinforcement Learning Introduction



HOCHSCHULE HEILBRONN

Prof. Dr.-Ing. Nicolaj Stache, Pascal Graf  
Heilbronn University of Applied Sciences

1. Deep Q-Learning
2. Task 1: 2D Pole Cart
3. DQN Improvements
4. Task 2: DQN Improvements
5. Policy Gradient
6. Actor Critic Methods
7. Task 3: Deep Deterministic Policy Gradient (DDPG)





# 1. Deep Q-Learning



HOCHSCHULE HEILBRONN

Prof. Dr.-Ing. Nicolaj Stache, Pascal Graf  
Heilbronn University of Applied Sciences

# From Tabular Methods to Deep Methods

250 states



<https://spiele.rtl.de/kartenspiele/black-jack.html>

$7.7 * 10^{45}$  states



<https://i.ytimg.com/vi/VHqCAaFXpbc/maxresdefault.jpg>

$10^{70802}$  states

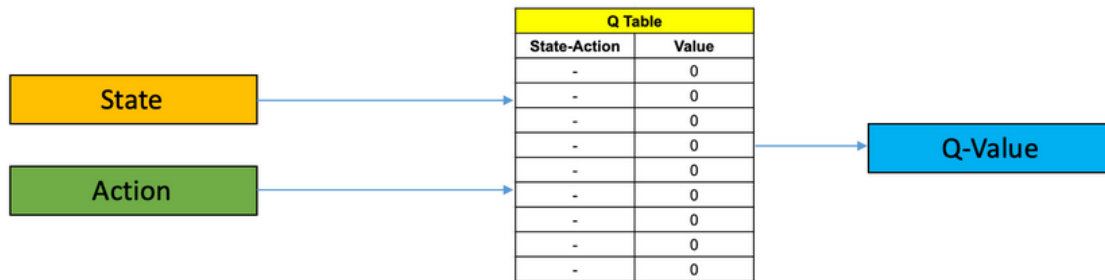


[https://www.retrogames.cz/play\\_222-Atari2600.php](https://www.retrogames.cz/play_222-Atari2600.php)


# From Tabular Methods to Deep Methods

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right] \quad (1)$$

$$y = R_{t+1} + \gamma \max_a Q(S_{t+1}, a) \quad L = (Q(S_t, A_t) - y)^2 \quad (2)$$



Q Learning



Deep Q Learning

<https://www.analyticsvidhya.com/blog/2019/04/introduction-deep-q-learning-python/>





1. Initialize  $Q(s, a)$  with some initial approximation.
2. By interacting with the environment, obtain the tuple  $(s, a, r, s')$ .
3. Calculate loss:  $\mathcal{L} = (Q(s, a) - r)^2$  if the episode has ended, or  
$$\mathcal{L} = \left( Q(s, a) - \left( r + \gamma \max_{a' \in A} Q(s', a') \right) \right)^2$$
 otherwise.
4. Update  $Q(s, a)$  using the **stochastic gradient descent (SGD)** algorithm, by minimizing the loss with respect to the model parameters.
5. Repeat from step 2 until converged.

<https://www.packtpub.com/data/deep-reinforcement-learning-hands-on-second-edition>

# Naive Algorithm

## PROBLEMS

- Exploration vs. Exploitation Dilemma:  
→ **Epsilon-Greedy Algorithm**
- Markov Property (**partially observable MDPs**)  
→ **State Stack**
- SGD optimization:  
*Training data needs to be **independent and identically distributed**.*  
→ **Replay Buffer**



# Naive Algorithm

## PROBLEMS

- Exploration vs. Exploitation Dilemma:  
    ➔ **Epsilon-Greedy Algorithm**
- Markov Property (**partially observable MDPs**)  
    ➔ **State Stack**
- SGD optimization:  
    *Training data needs to be **independent and identically distributed**.*  
    ➔ **Replay Buffer**
- Correlation between steps:  
    *We're training  $Q(s, a)$  via  $Q(s', a')$  (bootstrapping). When we perform an update of our NN's parameters to make  $Q(s, a)$ , we can indirectly alter the value produced for  $Q(s', a')$  and other states nearby.*  
    ➔ **Target Network**




$$L = \frac{1}{N} \sum_i (Q(s_i, a_i) - y_i)^2 \quad (3)$$



# Final Algorithm

The algorithm for DQN from the preceding papers has the following steps:

1. Initialize the parameters for  $Q(s, a)$  and  $\hat{Q}(s, a)$  with random weights,  $\varepsilon \leftarrow 1.0$ , and empty the replay buffer.
2. With probability  $\varepsilon$ , select a random action,  $a$ ; otherwise,  $a = \arg \max_a Q(s, a)$ .
3. Execute action  $a$  in an emulator and observe the reward,  $r$ , and the next state,  $s'$ .
4. Store transition  $(s, a, r, s')$  in the replay buffer.
5. Sample a random mini-batch of transitions from the replay buffer.
6. For every transition in the buffer, calculate target  $y = r$  if the episode has ended at this step, or  $y = r + \gamma \max_{a' \in \mathcal{A}} \hat{Q}(s', a')$  otherwise.
7. Calculate loss:  $\mathcal{L} = (Q(s, a) - y)^2$ .
8. Update  $Q(s, a)$  using the SGD algorithm by minimizing the loss in respect to the model parameters.
9. Every  $N$  steps, copy weights from  $Q$  to  $\hat{Q}$ .
10. Repeat from step 2 until converged.



<https://www.packtpub.com/data/deep-reinforcement-learning-hands-on-second-edition>

<https://www.analyticsvidhya.com/blog/2019/04/introduction-deep-q-learning-python/>



## 2. Task 1: 2D Pole Cart




HOCHSCHULE HEILBRONN

Prof. Dr.-Ing. Nicolaj Stache, Pascal Graf  
Heilbronn University of Applied Sciences

# 2D Pole Cart

## Task:

*Implement a Deep Q-Learning Algorithm (DQN) to balance a pole on a cart in two dimensions by moving the cart left and right. (Additionally, try to implement some of the improvement techniques discussed before and compare training performance.)*







### 3. DQN Improvements






HOCHSCHULE HEILBRONN

Prof. Dr.-Ing. Nicolaj Stache, Pascal Graf  
Heilbronn University of Applied Sciences

## MAXIMIZATION BIAS AND DOUBLE LEARNING

$$Target \leftarrow R_{t+1} + \gamma \max_a \hat{Q}(S_{t+1}, a) \quad (4)$$

$$Target \leftarrow R_{t+1} + \gamma \hat{Q}(S_{t+1}, \operatorname{argmax}_a Q(S_{t+1}, a)) \quad (5)$$



<http://www.andrew.cmu.edu/course/10-703/textbook/BartoSutton.pdf>

- Classical DQN achieves exploration with the hyperparameter epsilon, which is slowly decreased over time.
- Instead, the authors add a noise to the weights of fully-connected layers of the network and adjust the parameters of this noise during training using backpropagation.
- For every weight in a fully-connected layer, we have a random value that we draw from the normal distribution. Parameters of the noise  $\mu$  and  $\sigma$  are stored inside the layer and get trained using backpropagation, the same way that we train weights of the standard linear layer.

<https://arxiv.org/abs/1706.10295>


<https://www.packtpub.com/data/deep-reinforcement-learning-hands-on-second-edition>

# Improvements

## DUELING DQN

$$Q(S_t, A_t) = V(S_t) + A(S_t, A_t) \quad (6)$$

*Action-Value*      *State-Value*      *Advantage*



- The key motivation behind this architecture is that for some games, it is unnecessary to know the value of each action at every timestep.
- By explicitly separating two estimators, the dueling architecture can learn which states are (or are not) valuable, without having to learn the effect of each action for each state.
- Problem: The naive sum of the two is “unidentifiable,” in that given the  $Q$  value, we cannot recover the  $V$  and  $A$  uniquely.

$$Q(S_t, A_t) = V(S_t) + A(S_t, A_t) - \frac{1}{N} \sum_k A(S_t, k) \quad (7)$$

<https://arxiv.org/abs/1511.06581>

<https://www.packtpub.com/data/deep-reinforcement-learning-hands-on-second-edition>





## 4. Task 2: DQN Improvements



HOCHSCHULE HEILBRONN

Prof. Dr.-Ing. Nicolaj Stache, Pascal Graf  
Heilbronn University of Applied Sciences





## 5. Policy Gradient



HOCHSCHULE HEILBRONN

Prof. Dr.-Ing. Nicolaj Stache, Pascal Graf  
Heilbronn University of Applied Sciences

# Policy Gradient

## VALUE VS. POLICY GRADIENT


### Previous Algorithms:

- Estimate **value**  $V(s)$  or  $Q(s, a)$
- Take the action with the highest estimated value in every state

$$\pi(s) = \operatorname{argmax}_a Q(s, a) \quad (8)$$

### Why Policy?

1. Environments with lots of actions or a **continuous action space**



<https://www.packtpub.com/data/deep-reinforcement-learning-hands-on-second-edition>

# Policy Gradient

## VALUE VS. POLICY GRADIENT


### Previous Algorithms:

- Estimate **value**  $V(s)$  or  $Q(s, a)$
- Take the action with the highest estimated value in every state

$$\pi(s) = \operatorname{argmax}_a Q(s, a) \quad (8)$$

### Why Policy?

1. Environments with lots of actions or a **continuous action space**
2. Environments with **stochasticity** in them



<https://www.packtpub.com/data/deep-reinforcement-learning-hands-on-second-edition>



# Policy Gradient

## VALUE VS. POLICY GRADIENT


### Previous Algorithms:

- Estimate **value**  $V(s)$  or  $Q(s, a)$
- Take the action with the highest estimated value in every state

$$\pi(s) = \operatorname{argmax}_a Q(s, a) \quad (8)$$

### Why Policy?

1. Environments with lots of actions or a **continuous action space**
2. Environments with **stochasticity** in them
3. Enables smooth representation



<https://www.packtpub.com/data/deep-reinforcement-learning-hands-on-second-edition>

# Policy Gradient

## POLICY GRADIENT

### Key Idea

Push up the probabilities of actions that lead to higher return and push down the probabilities of actions that lead to lower return, until you arrive at the optimal policy.

### Policy performance:

$$J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta}[G(\tau)] \quad (9)$$

### Weight update for gradient ascent:

$$\theta_{k+1} = \theta_k + \alpha \underbrace{\nabla_\theta J(\pi_\theta)}_{\text{Policy Gradient}} \quad (10)$$

$$\nabla_\theta J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t | s_t) G(\tau) \right] \approx \frac{1}{|D|} \sum_{\tau \in D} \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t | s_t) G(\tau) \quad (11)$$

[https://medium.com/@jonathan\\_hui/rl-policy-gradients-explained-9b13b688b146](https://medium.com/@jonathan_hui/rl-policy-gradients-explained-9b13b688b146)

[https://spinningup.openai.com/en/latest/spinningup/rl\\_intro3.html](https://spinningup.openai.com/en/latest/spinningup/rl_intro3.html)

<https://www.packtpub.com/data/deep-reinforcement-learning-hands-on-second-edition>

## Pseudocode

1. Initialize the network with random weights.
2. Play  $D$  full episodes, saving their  $(s, a, r, s')$  transitions.
3. For every step  $t$  of every trajectory  $\tau$ , calculate the discounted total reward for subsequent steps  $G_{\tau,t} = \sum_{i=0}^{\infty} \gamma^i r_{t+i}$ .
4. Calculate the loss function for all transitions.

$$L = -\frac{1}{|D|} \sum_{\tau \in D} \sum_{t=0}^T \log \pi_{\theta}(a_t | s_t) G(\tau) \quad (12)$$

5. Perform SGD update of weights minimizing the loss.
6. Repeat from step 2 until converged.

$$\nabla_{\theta} J(\pi_{\theta}) \approx \frac{1}{|D|} \sum_{\tau \in D} \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau) \quad (11)$$

[https://medium.com/@jonathan\\_hui/rl-policy-gradients-explained-9b13b688b146](https://medium.com/@jonathan_hui/rl-policy-gradients-explained-9b13b688b146)

[https://spinningup.openai.com/en/latest/spinningup/rl\\_intro3.html](https://spinningup.openai.com/en/latest/spinningup/rl_intro3.html)

<https://www.packtpub.com/data/deep-reinforcement-learning-hands-on-second-edition>



## 6. Actor Critic Methods



HOCHSCHULE HEILBRONN

Prof. Dr.-Ing. Nicolaj Stache, Pascal Graf  
Heilbronn University of Applied Sciences



## Shortcomings of Policy Gradient:

- Whole trajectories needed
- High variability in log probabilities and cumulative rewards:


→ **High variance gradients**

- Trajectories with cumulative reward of zero:

→ **Zero gradients**

- On-Policy

$$\nabla_{\theta} J(\pi_{\theta}) \approx \frac{1}{|D|} \sum_{\tau \in D} \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau) \quad (11)$$



$R_t = -0.1$   
on all other transitions

<https://towardsdatascience.com/understanding-actor-critic-methods-931b97b6df3f>

[https://medium.com/@jonathan\\_hui/rl-policy-gradients-explained-9b13b688b146](https://medium.com/@jonathan_hui/rl-policy-gradients-explained-9b13b688b146)

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau) \right] \quad (11)$$

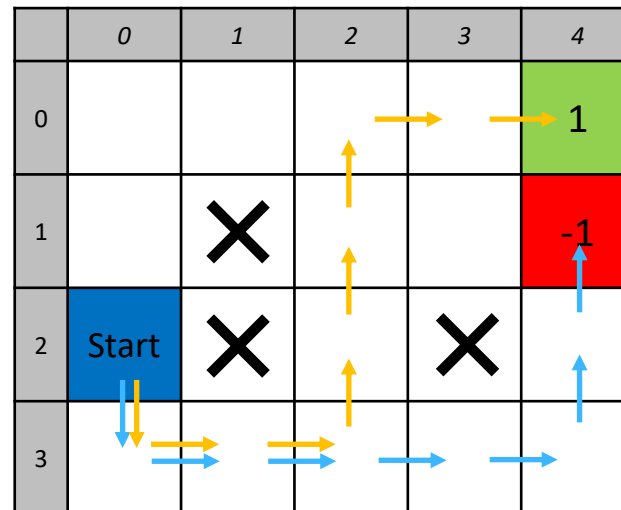
$$= \mathbb{E}_{s_0, a_0, \dots, s_t, a_t} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right] \mathbb{E}_{r_{t+1}, \dots, r_T} [G(\tau)]$$

$$= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) Q_{\omega}(s_t, a_t) \right] \quad (12)$$

<https://towardsdatascience.com/understanding-actor-critic-methods-931b97b6df3f>

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau) \right] \quad (11)$$

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) Q_{\omega}(s_t, a_t) \right] \quad (12)$$

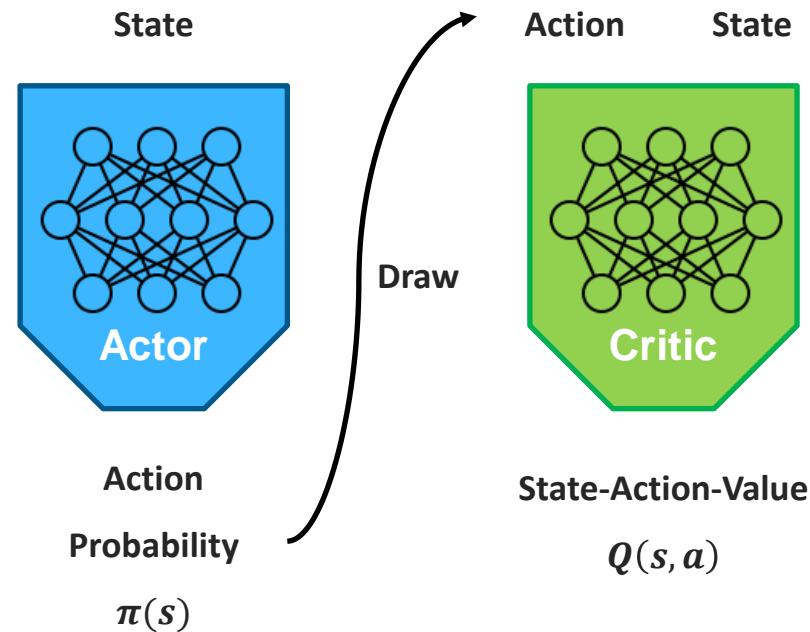


$R_t = -0.1$   
on all other transitions

<https://towardsdatascience.com/understanding-actor-critic-methods-931b97b6df3f>

# Q Actor Critic

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) Q_{\omega}(s_t, a_t) \right] \quad (12)$$



<https://towardsdatascience.com/understanding-actor-critic-methods-931b97b6df3f>



$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) Q_{\omega}(s_t, a_t) \right] \quad (12)$$

---

**Algorithm 1** Q Actor Critic

---

Initialize parameters  $s, \theta, w$  and learning rates  $\alpha_{\theta}, \alpha_w$ ; sample  $a \sim \pi_{\theta}(a|s)$ .  
**for**  $t = 1 \dots T$ : **do**  
    Sample reward  $r_t \sim R(s, a)$  and next state  $s' \sim P(s'|s, a)$   
    Then sample the next action  $a' \sim \pi_{\theta}(a'|s')$   
    Update the policy parameters:  $\theta \leftarrow \theta + \alpha_{\theta} Q_w(s, a) \nabla_{\theta} \log \pi_{\theta}(a|s)$ ; Compute the correction (TD error) for action-value at time t:  
         $\delta_t = r_t + \gamma Q_w(s', a') - Q_w(s, a)$   
    and use it to update the parameters of Q function:  
         $w \leftarrow w + \alpha_w \delta_t \nabla_w Q_w(s, a)$   
    Move to  $a \leftarrow a'$  and  $s \leftarrow s'$   
**end for**

---

## Benefits of Q Actor Critic:

- Updates after one step of playing
- Lower variance in policy gradients

## Shortcomings of Q Actor Critic:

- On-Policy algorithm


<https://towardsdatascience.com/understanding-actor-critic-methods-931b97b6df3f>

# Deep Deterministic Policy Gradient (DDPG)

## Ideas of DDPG:

- Continuous actions instead of probability distribution
- Off-Policy Learning (→ Replay Buffer)
- Online learning

State



Action

State



State-Action-Value

$$Q(s, a)$$

<https://towardsdatascience.com/deep-deterministic-and-twin-delayed-deep-deterministic-policy-gradient-with-tensorflow-2-x-43517b0e0185>

<https://towardsdatascience.com/deep-deterministic-policy-gradients-explained-2d94655a9b7b>

# Deep Deterministic Policy Gradient (DDPG)

## Actor Critic Policy Gradient:

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) Q_{\omega}(s_t, a_t) \right] \quad (12)$$


## DDPG:

$$J(\mu_{\theta}) = \mathbb{E}[Q_{\omega}(s_t, \mu_{\theta}(s))] \quad (13)$$

$$\nabla_{\theta} J(\mu_{\theta}) = \mathbb{E}[\nabla_{\mu_{\theta}} Q_{\omega}(s_t, \mu_{\theta}(s_t)) \nabla_{\theta} \mu_{\theta}(s_t)]$$

$$\approx \frac{1}{N} \sum_i \nabla_{\mu_{\theta}} Q_{\omega}(s_i, \mu_{\theta}(s_i)) \nabla_{\theta} \mu_{\theta}(s_i) \quad (14)$$

State



Action

State



State-Action-Value

$Q(s, a)$

<https://towardsdatascience.com/deep-deterministic-and-twin-delayed-deep-deterministic-policy-gradient-with-tensorflow-2-x-43517b0e0185>

<https://towardsdatascience.com/deep-deterministic-policy-gradients-explained-2d94655a9b7b>



# Deep Deterministic Policy Gradient (DDPG)

## DDPG Actor Update:

$$\nabla_{\theta} J(\mu_{\theta}) \approx \frac{1}{N} \sum_i \nabla_{\mu_{\theta}} Q_{\omega}(s_i, \mu_{\theta}(s_i)) \nabla_{\theta} \mu_{\theta}(s_i) \quad (14)$$


## Q Learning Target:

$$y_i = r_i + \gamma \max_{a'} \hat{Q}(s'_i, a') \quad L = \frac{1}{N} \sum_i (Q(s_i, a_i) - y_i)^2 \quad (3)$$

## DDPG Critic Target:

$$y_i = r_i + \gamma \hat{Q}(s'_i, \hat{\mu}(s'_i)) \quad L_Q = \frac{1}{N} \sum_i (Q(s_i, a_i) - y_i)^2 \quad (15)$$

State



Action

State



State-Action-Value

$Q(s, a)$

<https://towardsdatascience.com/deep-deterministic-and-twin-delayed-deep-deterministic-policy-gradient-with-tensorflow-2-x-43517b0e0185>

<https://towardsdatascience.com/deep-deterministic-policy-gradients-explained-2d94655a9b7b>

---

**Algorithm 1** DDPG algorithm

---

Randomly initialize critic network  $Q(s, a|\theta^Q)$  and actor  $\mu(s|\theta^\mu)$  with weights  $\theta^Q$  and  $\theta^\mu$ .  
Initialize target network  $Q'$  and  $\mu'$  with weights  $\theta^{Q'} \leftarrow \theta^Q$ ,  $\theta^{\mu'} \leftarrow \theta^\mu$   
Initialize replay buffer  $R$   
**for** episode = 1, M **do**  
    Initialize a random process  $\mathcal{N}$  for action exploration  
    Receive initial observation state  $s_1$   
    **for** t = 1, T **do**  
        Select action  $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$  according to the current policy and exploration noise  
        Execute action  $a_t$  and observe reward  $r_t$  and observe new state  $s_{t+1}$   
        Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $R$   
        Sample a random minibatch of  $N$  transitions  $(s_i, a_i, r_i, s_{i+1})$  from  $R$   
        Set  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$   
        Update critic by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$   
        Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

**end for**  
**end for**

---



## 7. Task 3: Deep Deterministic Policy Gradient (DDPG)




HOCHSCHULE HEILBRONN

Prof. Dr.-Ing. Nicolaj Stache, Pascal Graf  
Heilbronn University of Applied Sciences



# Further Readings



[https://medium.com/@jonathan\\_hui/rl-policy-gradients-explained-9b13b688b146](https://medium.com/@jonathan_hui/rl-policy-gradients-explained-9b13b688b146)

[https://spinningup.openai.com/en/latest/spinningup/rl\\_intro3.html](https://spinningup.openai.com/en/latest/spinningup/rl_intro3.html)

<https://www.packtpub.com/data/deep-reinforcement-learning-hands-on-second-edition>

- **TD3 (TWIN DELAYED DDPG)**
- **PPO (PROXIMAL POLICY OPTIMIZATION)**
- **SAC (SOFT ACTOR-CRITIC)**